

J-space Phase 4 — development synthesis

OLMo trajectory, coordinate validity, and Qwen fit-size gate

State through the registered OLMo trajectory and Qwen draw-A $n = 120$ structural comparison · 2026-07-31

Scope. This is a living **development only** report over known Phase 3 banks. It localizes trajectory patterns and audits lens-frame dependence; it cannot establish a binary lineage claim. No Phase 4 confirmatory or replication cell is licensed before PI sign-off and a tagged preregistration freeze.

1 Compute and provenance boundary

Every model producer hard-failed unless CUDA was visible in the same process, executed a finite FP16 CUDA matrix multiply before model load, and recorded the device in its result envelope. Jobs ran on an NVIDIA RTX PRO 6000 Blackwell Server Edition; the final grid allocated about 81.1 GB VRAM. Model loading, generation, interventions, and scoring never used CPU. CPU was limited to tests, hashes, statistics, figures, and document compilation.

Phase 3 enters through the immutable `p4-import-phase3-release-v1`. Native evidence is event-sourced; registered artifacts are never overwritten.

2 Prospective capability gate

OLMo-3 32B Think produced 972 G5 rows, 324 facts, and 72 families under prospective prefix-disjoint accepted-alias scoring.

view	overall	Bank F	Bank S	direct	composed	bridge supplied
capability	.6430	.4935	.8972	.6574	.4846	.7870

The intervention cohort was fixed before outcomes by requiring both direct and composed generation capability within a fact: 41/204 Bank-F facts and 85/120 Bank-S facts, or 126 facts / 252 items.

Evidence: `p4-g5-bank-olmo3-think-dev-v2`

The first attempt stopped after accent folding made `Río` and `Rio` ambiguous. The repaired selector prefers exact spelling before a unique normalized fallback; the partial attempt remains preserved.


3 OLMo-3 32B base point

The exact pinned base revision `c2b61dae89a1ad10e4ad5653d0e46b590902607b` produced 972 G5 rows. Overall capability was .6101 (Bank F .4395; Bank S .9000). Requiring both direct and composed capability fixed 23 Bank-F and 88 Bank-S facts, or 111 facts / 222 items.

Evidence: `p4-g5-bank-olmo3-base-dev-v1`

The seven-condition grid had maximum baseline replay drift 6.71×10^{-7} nats, zero span-safe selected/protected overlap, 100% matched-rank agreement, and maximum energy relative error 0.000260.

Evidence: `p4-lineage-grid-olmo3-base-dev-v1`



cell	J-specific estimate (nats)	family-bootstrap 95% interval
Bank F direct	+0.0245	[−0.0330, +0.0792]
Bank F composed	+0.0145	[−0.1226, +0.1881]
Bank S direct	+0.0005	[−0.0487, +0.0481]
Bank S composed	+0.0021	[−0.0309, +0.0385]

All four primary effects are near zero. Bank-S composed-minus-direct is $+0.0016$ $[-0.0417, +0.0403]$. The experiment is nevertheless sensitive: overall label-protected J-specific is $+0.1231$ $[+0.0667, +0.1802]$, while mechanics-random and logit-protected effects are -0.1048 and -0.2023 , with intervals below zero.


Evidence: [p4-lineage-analysis-olmo3-base-dev-v1](#)

4 Think own-lens development point

The seven-condition grid includes baseline, span-safe J, exact rank/energy control, label-protected J, protected-energy control, mechanics random, and logit-space label protection. Baseline replay drift is at most 7.15×10^{-7} nats. Rank agreement is 100%, energy relative error is at most 0.000230, and span-safe selected/protected overlap is zero.

Evidence: [p4-lineage-grid-olmo3-think-dev-v1](#)

OLMo-3 32B Think: Phase 4 development trajectory point



cell	J-specific estimate (nats)	family-bootstrap 95% interval
Bank F direct	+0.0574	$[-0.0442, +0.1785]$
Bank F composed	+0.0818	$[-0.0223, +0.1878]$
Bank S direct	-0.1277	$[-0.2072, -0.0494]$
Bank S composed	-0.0553	$[-0.0982, -0.0187]$

Bank-S composed-minus-direct specificity is $+0.0724 [-0.0109, +0.1584]$. It is compatible with the unusual 3.1 Think tendency beginning by 3.0, but is not resolved and must await the controlled Bank-W redundancy/derivation study.

Evidence: [p4-lineage-analysis-olmo3-think-dev-v1](#)

5 Common-base-lens repair audit

The common coordinate is the frozen OLMo base lens (92f32e...b696). V1 exposed rank-limited protected-energy sites that omitted one requested component without being marked clamped. V2 repaired the metadata but revealed that the evidence ID itself changed all matched/random seeds.

V3 freezes the scientific-seed namespace to v1. All seven aggregate outcome columns, every per-alias score, and every condition order match v1 exactly (maximum difference 0). The fix reclassifies 296 summaries and reduces the erroneous maximum unclamped energy error from 0.511283 to 0.000239 while retaining 100% rank agreement. V3 supersedes both withdrawn diagnostic attempts and remains valid standalone evidence.

Evidence: [p4-lineage-grid-olmo3-think-common-base-lens-dev-v3](#)


V4 reruns the common lens under the own-grid seed namespace. Its 252 rows form 126 exact direct/composed pairs; condition order matches the own grid item-for-item, between-frame baseline drift is zero, matched rank agreement is 100%, maximum energy error is 0.000205, and span-safe overlap is zero. Deterministic J arms match v3 exactly.

Evidence: [p4-lineage-grid-olmo3-think-common-base-lens-dev-v4](#)

6 Own/common lens comparison: corrected

The v1 paired analysis remains **withdrawn** because its frames used different matched-control RNG streams. V2 shares the exact namespace and condition order. Its full 100,000-draw payload and all bootstrap hashes reproduce exactly. The hardened producer refuses unequal namespaces.

Evidence: p4-lens-frame-analysis-olmo3-think-dev-v2



Same 126 paired facts; known development banks; family-resampling 95% intervals; not confirmatory evidence.

cell	own lens	common lens	common - own (95% interval)
F direct	+0.0574	+0.0010	-0.0563 [-0.1324, +0.0049]
F composed	+0.0818	-0.0164	-0.0982 [-0.2005, +0.0058]
S direct	-0.1277	-0.0974	+0.0303 [-0.0165, +0.0813]
S composed	-0.0553	-0.0419	+0.0134 [-0.0100, +0.0408]

No frame-delta interval excludes zero. Bank-S direct and composed effects have intervals below zero in each tested frame; Bank-F effects remain imprecise. Item and family frame correlations are .7557 and .7254; mean absolute item difference is .1025 nats. The RNG correction leaves J deltas unchanged, moves Bank-F specificity deltas 0.0413 and 0.0492 nats toward zero through the controls, and raises family correlation from .4946 to .7254.

7 OLMo-3.1 Think capability and own frame


The exact pinned 3.1 Think revision 832c3f54...7840544 produced 972 G5 rows. Boundary-safe prefix capability was .6327 overall (Bank F .4837; Bank S .8861), with direct .6420, composed .4753, and bridge-supplied .7809. Requiring both direct and composed capability fixed 38 Bank-F and 84 Bank-S facts, or 122 facts / 244 items.

Evidence: p4-g5-bank-olmo31-think-dev-v1

This rate is below Phase 3's historical .7346 because that cohort allowed an answer anywhere in the continuation. All 972 generations are byte-identical across phases, and every Phase 4 flag equals the later Phase 3 boundary-safe prefix audit.

The own-lens grid has maximum G5 replay drift 8.64×10^{-7} nats, zero span-safe overlap, 100% rank agreement over 564 matched summaries, and maximum energy error 0.000222.

Evidence: p4-lineage-grid-olmo31-think-dev-v1



Known development banks; family-resampling 95% intervals; not confirmatory evidence.

cell	J-specific estimate (nats)	family-bootstrap 95% interval
Bank F direct	-0.0067	[-0.1162, +0.1057]
Bank F composed	-0.0416	[-0.2099, +0.1167]
Bank S direct	-0.1674	[-0.2477, -0.0942]
Bank S composed	-0.0491	[-0.0891, -0.0165]


Bank-S composed-minus-direct is $+0.1183 [+0.0514, +0.1879]$ (descriptive exact sign-flip $p = .0031$). Analysis v1 has an identical numerical payload but is withdrawn for a crowded figure footer; v2 fixes the reusable layout.

Evidence: [p4-lineage-analysis-olmo31-think-dev-v2](#)

8 OLMo-3.1 Think paired coordinate comparison

The common grid uses the frozen base lens under the exact own-grid seed namespace. All 244 item IDs and condition orders pair; between-frame baseline drift is zero. Its maximum matched energy error is 0.000256, with zero span-safe overlap and 100% rank agreement. Mechanics-random and logit-protected outcomes are bit-identical between frames.

Evidence: [p4-lineage-grid-olmo31-think-common-base-lens-dev-v1](#)



Same 122 paired facts; known development banks; family-resampling 95% intervals; not confirmatory evidence.

cell	own lens	common lens	common - own (95% interval)
F direct	-0.0067	-0.0757	-0.0690 [-0.1413, +0.0025]
F composed	-0.0416	-0.1314	-0.0898 [-0.1677, -0.0096]
S direct	-0.1674	-0.1547	+0.0127 [-0.0378, +0.0669]
S composed	-0.0491	-0.0381	+0.0110 [-0.0196, +0.0406]

Only the Bank-F composed frame-delta interval excludes zero: evidence of coordinate sensitivity for that known-bank cell, although the effect itself remains imprecise in each frame. Bank-S direct is negative in both frames. The Bank-S composition contrast is positive in both frames: own +0.1183 [+0.0517, +0.1876], common +0.1166 [+0.0279, +0.2080]; paired delta -0.0017 [-0.0551, +0.0504]. Item/family correlations are .7679/.8386 and mean absolute item shift is .1133 nats. All 42 bootstrap and 6 exact sign-flip distributions reproduce exactly.

Evidence: [p4-lens-frame-analysis-olmo31-think-dev-v1](#)


9 OLMo-3.1 Instruct capability and own frame

The exact pinned 3.1 Instruct revision `ac0587e4...bc940dae` produced 972 G5 rows. Boundary-safe prefix capability was .6245 overall (Bank F .4673; Bank S .8917), with direct .6667, composed .4383, and bridge-supplied .7685. Requiring both direct and composed capability fixed 32 Bank-F and 86 Bank-S facts, or 118 facts / 236 items. All 972 generations match Phase 3 byte-for-byte and every flag matches its later boundary-safe audit.

Evidence: [p4-g5-bank-olmo31-instruct-dev-v1](#)

The own grid has maximum G5 replay drift 7.75×10^{-7} nats, zero span-safe overlap, 100% rank agreement over 528 matched summaries, and maximum energy error 0.000228.

Evidence: [p4-lineage-grid-olmo31-instruct-dev-v1](#)



Known development banks; family-resampling 95% intervals; not confirmatory evidence.

cell	J-specific estimate (nats)	family-bootstrap 95% interval
Bank F direct	+0.0432	$[-0.0259, +0.1181]$
Bank F composed	-0.0662	$[-0.4037, +0.1921]$
Bank S direct	-0.0224	$[-0.0656, +0.0179]$
Bank S composed	-0.0177	$[-0.0755, +0.0325]$


Bank-S composed-minus-direct is $+0.0047 [-0.0411, +0.0501]$. All four primary intervals include zero, unlike the 3.1 Think Bank-S cells. All 22 bootstrap and 6 exact sign-flip distributions reproduce exactly.

Evidence: `p4-lineage-analysis-olmo31-instruct-dev-v1`

10 OLMo-3.1 Instruct paired coordinate comparison

The own/common grids have identical 236-item manifests, condition orders, and scientific RNG namespaces. Between-frame baseline drift is zero; each frame replays G5 within 7.75×10^{-7} nats. Span-safe overlap is zero and rank agreement is 100%. Maximum energy error is 0.000228 own and 0.000212 common. Mechanics-random and logit-protected outcomes are bit-identical.

Evidence: `p4-lineage-grid-olmo31-instruct-common-base-lens-dev-v1`



Same 118 paired facts; known development banks; family-resampling 95% intervals; not confirmatory evidence.

cell	own lens	common lens	common – own (95% interval)
F direct	+0.0432	–0.0111	–0.0542 [–0.0970, –0.0194]
F composed	–0.0662	+0.0271	+0.0934 [–0.0312, +0.2488]
S direct	–0.0224	–0.0379	–0.0155 [–0.0603, +0.0258]
S composed	–0.0177	–0.0045	+0.0131 [–0.0240, +0.0558]

Bank-F direct has a negative frame delta with interval below zero (descriptive exact sign-flip $p = .0057$). Bank-F composition changes from own $-0.1094 [-0.4439, +0.1404]$ to common $+0.0382 [-0.1652, +0.2172]$; paired delta $+0.1476 [+0.0257, +0.2978]$ ($p = .0363$). All Bank-S effects, composition contrasts, and frame deltas include zero. Item/family correlations are .7523/.8521 and mean absolute item shift is .0819 nats. All 42 bootstrap and 6 exact sign-flip distributions reproduce exactly.


Evidence: `p4-lens-frame-analysis-olmo31-instruct-dev-v1`

11 Registered four-checkpoint trajectory

The synthesis reads all immutable own/common grids under one stable family-bootstrap schedule. It requires live evidence, a shared bank/scoring contract, exact own/common item pairing, and identical scientific RNG namespaces and condition orders. All 48 table summaries and all 42 unique bootstrap distributions reconstruct exactly.

Evidence: `p4-lineage-trajectory-analysis-olmo-dev-v1`

OLMo 32B development trajectory in two coordinate frames



checkpoint	S direct own	S direct common	S composition own	S composition common
Base	+0.0005	+0.0005	+0.0016	+0.0016
3.0 Think	-0.1277	-0.0974	+0.0724	+0.0556
3.1 Think	-0.1674	-0.1547	+0.1183	+0.1166
3.1 Instruct sibling	-0.0224	-0.0379	+0.0047	+0.0333

The figure connects only Base → 3.0 Think → 3.1 Think; squares mark the sibling Instruct endpoint, not a temporal continuation. In both frames, Bank-S direct becomes negative by 3.0 and more negative by 3.1, while Bank-S composition rises from near zero at base to about +0.117 at 3.1 Think. Both patterns collapse toward zero at sibling Instruct. Checkpoint-specific capability cohorts make this an unpaired development localization, not a causal training or mode contrast.

12 Qwen fit-size gate: first nested lens

The leakage-safe draw-A corpus has exact nested prefixes at $n = 120/250/500/1000$. Its first 120 records reproduce the historical fit prefix byte-for-byte; all 80 historical evaluation spares are excluded. Independent draw B is disjoint and also excludes every spare.

Evidence: [p4-qwen-lens-corpora-dev-v1](#)

The first all-layer lens fits L0–L62 to target L63 at width 5120, using 128 tokens and a fixed 16-token skip. Every model operation ran on the RTX PRO 6000 under pinned fused FLA kernels; all 48 linear-attention blocks passed the binding gate. Forty three-prompt checkpoints were hashed and mirrored to Drive. The 6.606 GB final checkpoint has SHA-256 061574f9...cdb673; the registered 3.303 GB fp16 lens is 82af4cc7...c32b1b. Full runtime was 21,622.7 seconds with 62.83 GB peak allocated VRAM.

Evidence: [p4-qwen-lens-fit-drawA-n120-dev-v1](#)

Prompt 112 was a pronounced but finite Jacobian-norm outlier (159.952). It remains in the frozen estimator; no post-hoc trimming is permitted.


13 Qwen structural stability is depth-dependent

The first diagnostic compares draw-A $n = 120$ with the published $n = 1000$ lens. It is a recipe/corpus-transfer comparison, not the pending same-corpus convergence test. It uses all 63 maps and all 5,120 Jacobian rows, 4,096

fixed token directions, centered linear CKA on a fixed 1,024-token subset, and 256 fixed transport probes.

Evidence: p4-qwen-lens-structural-stability-drawA-n120-vs-published-n1000-dev-v1

Qwen3.6-27B lens structure: draw-A n=120 vs published n=1000



Development recipe/corpus-transfer diagnostic; not the pending same-corpus nested convergence test.

layer view	matrix cosine	token median	token q05	CKA	relative $\ \Delta\ _F$
L0	.5411	.5297	.2760	.4518	1.1086
L24	.8855	.8947	.8145	.8667	.4711
L32	.9043	.9220	.8263	.8409	.4269
L40	.9605	.9677	.9312	.9433	.2800
L62	.9997	.9997	.9994	.9997	.0234
layer median	.8969	.9138	.8187	.8652	.4422

Token median first remains at least .90 from L26, matrix cosine from L32, CKA from L33, and token q05 from L35. Thus $n = 120$ nearly matches the published lens late but is not uniformly stable in the lower/middle assay layers. The result triggers the same-corpus $n = 250$ fit. It does not yet satisfy selected-ID, span, occupancy, centered capacity, G4, or causal/bridge stability gates. Both envelopes reconstruct; all 63 table rows are finite; the live verifier passes 25 events / 97 outputs / zero failures.

14 Interpretation and next boundary

1. The base checkpoint is near zero in all four primary cells. Bank-S direct and composed specificity is negative by 3.0 Think and is negative with intervals below zero in both seed-paired frames. The paired frame deltas include zero: development evidence of two-frame robustness, not a confirmatory invariance claim.
2. Because capability cohorts are checkpoint-specific, this localizes a base-to-Think trajectory change but is not a paired causal estimate.
3. At 3.1 Think, Bank-S direct remains negative in both frames. Bank-S composed is negative in the own frame but imprecise in the common frame. The positive composition contrast is precise in both frames and stable under the lens change.
4. At sibling 3.1 Instruct, every Bank-S specificity and composition interval includes zero in both frames. The Think pattern does not transfer under these development cohorts; differing capability cohorts and post-training objectives preclude a paired causal mode claim.
5. Coordinate choice affects Bank-F localization: 3.1 Think has a negative composed frame delta; 3.1 Instruct has a negative direct delta and positive composition delta with intervals excluding zero.
6. The 3.1 composition result is a development finding, not a mechanism; controlled Bank-W load/redundancy and internal-derivation axes remain necessary.
7. None of these known-bank estimates licenses “lineage present” or “lineage absent.”

A5 and Qwen draw-A $n = 120$ are banked; depth-dependent drift triggers same-corpus $n = 250$. Next: fit-size/model-backed stability, Bank-W load/redundancy, internal derivation, and an intermediate-checkpoint audit.