Lab 38 Phase 1, part 1 — instrument build and development pilot

Stated vs revealed preference: counterbalanced action-binding choice on `Olmo-3-7B-Instruct`

Branch `interp_preference_phase1` — 2026-08-07 — freeze-review input, pre-registration candidate attached

Scope. This is a living development only report. Every number below is development-tier evidence produced before the freeze; nothing here is a preference result, and no frozen-partition outcome was opened. The claim ceiling applies verbatim: the campaign studies *functional* choice, report, and their coupling — never wants, welfare, suffering, consent, experience, moral patienthood, introspective truth, or a preference workspace. This handout is a freeze-review input, not a review signature, PI approval, freeze commit, or publication claim.

1 What was built

A governed campaign under `interpretability/preference/` (plan `preference_1_1.md` + binding addendum; both hash-pinned at intake). The two draft generators and the design transcript were lost before intake (`missing_at_intake`, PI-confirmed); the `lab38_v2_phase1` instrument was built from scratch with every plan-P0 repair native:

- **Bank**: 2,320 rows — 12 arbitrary (AR), 6 positive-control (PC: 2 quality / 2 social / 2 safety), 2 null-control (NC, verbatim-identical options) scenarios × 5 incidentals × full counterbalance (2 orders × 2 display-label sets × 2 response-code maps × 2 consequence frames), plus a report-only (RO) twin for every AR/PC cell. Splits 3 train / 1 validation / 1 holdout per scenario; construct axes `naming_convention`×3 and `execution_mode`×2 make construct-level claims reachable. Scientific identity: every behavior-relevant field is canonically hashed into the item id. Regeneration is byte-deterministic.
  Evidence: `pref1-schema-audit-v1`
- **Response-code contract**: opaque consonant-consonant-digit codes audited on the pinned tokenizer per the addendum-E procedure; AR pair `KP4`/`PK7` (neutral-prior gap 0.0031 nats), RO pair `VM2`/`GS2` (0.0275 nats); disjoint alphabets, distinct first tokens, no prefix relations; codes listed in display order (E12).
  Evidence: `pref1-codebook-final-v1`
- **Strict parser**: exact-code-or-invalid, never guesses; 12/12 adversarial cases behave as specified. **Binding**: enacted valid choices execute exactly the parsed branch (4 scenarios carry model microtasks with deterministic validators); hypothetical, report-only, and invalid rows never execute (E10); wrong-branch executions are a hard stop.
- **Runner**: canonical-order batches, immutable per-item JSONL with fsync, atomic resume cursor that refuses config changes, interrupted-resume byte-parity proven, Drive mirroring at batch boundaries. **Equality review**: two blinded provisional agent passes (disagreements preserved), both pre-model; PI ratings gate the freeze. 46 unit + synthetic-analysis tests green.
  Evidence: `pref1-gate-g2-instrument-v1`

2 Instrument finding: bf16 batched scoring is not batch-invariant

On the 7B in bf16, batched teacher-forced margins differed from single-row scoring by up to **0.250 nats** (kernel/shape numerics; strict generations were identical). Margins are therefore scored **single-row**: exact, batch-shape-free, and resume-invariant by construction. Per-run hard gates: single-row replay determinism and batched=single generations; the batched delta is recorded as an informational diagnostic. This is one of the most reusable numbers this part produced — any lab scoring small margins from batched bf16 forwards should check it.

Evidence: `diagnostics/batch_invariance.json`

3 Development pilot (696 rows, dev subset)

Dev subset = train incidentals × both orders × letter labels × both code maps × both frames. (The first pilot slice used one order only; NC immediately exposed the aliasing — a deterministic tie-break registered as a spurious |effect| = 0.5 — so the slice was widened before the freeze. Dev exists to catch exactly this.) Runtime: 10.7 rows/s measured; full frozen bank projects to ∼9 GPU-minutes.

check	observed	gate
strict valid-parse rate (all 696)	0.9986	—
PC strict parse rate	1.000	≥ 0.98 PASS
PC expected-content rate	1.000	≥ 0.85 PASS
PC per-scenario minimum	1.000	≥ 0.75 PASS
binding execution (valid enacted)	1.000	$= 1.00$ PASS
wrong-branch executions	0	$= 0$ PASS
NC false-positive floor (p95, dev n)	0.208	informational
microtask follow-through pass	0.875	exploratory
AR↔RO agreement (enacted)	0.693	descriptive

The single invalid generation in 696 rows was `PK4` — a blend of the two AR codes; the strict parser correctly refused it and no branch executed.

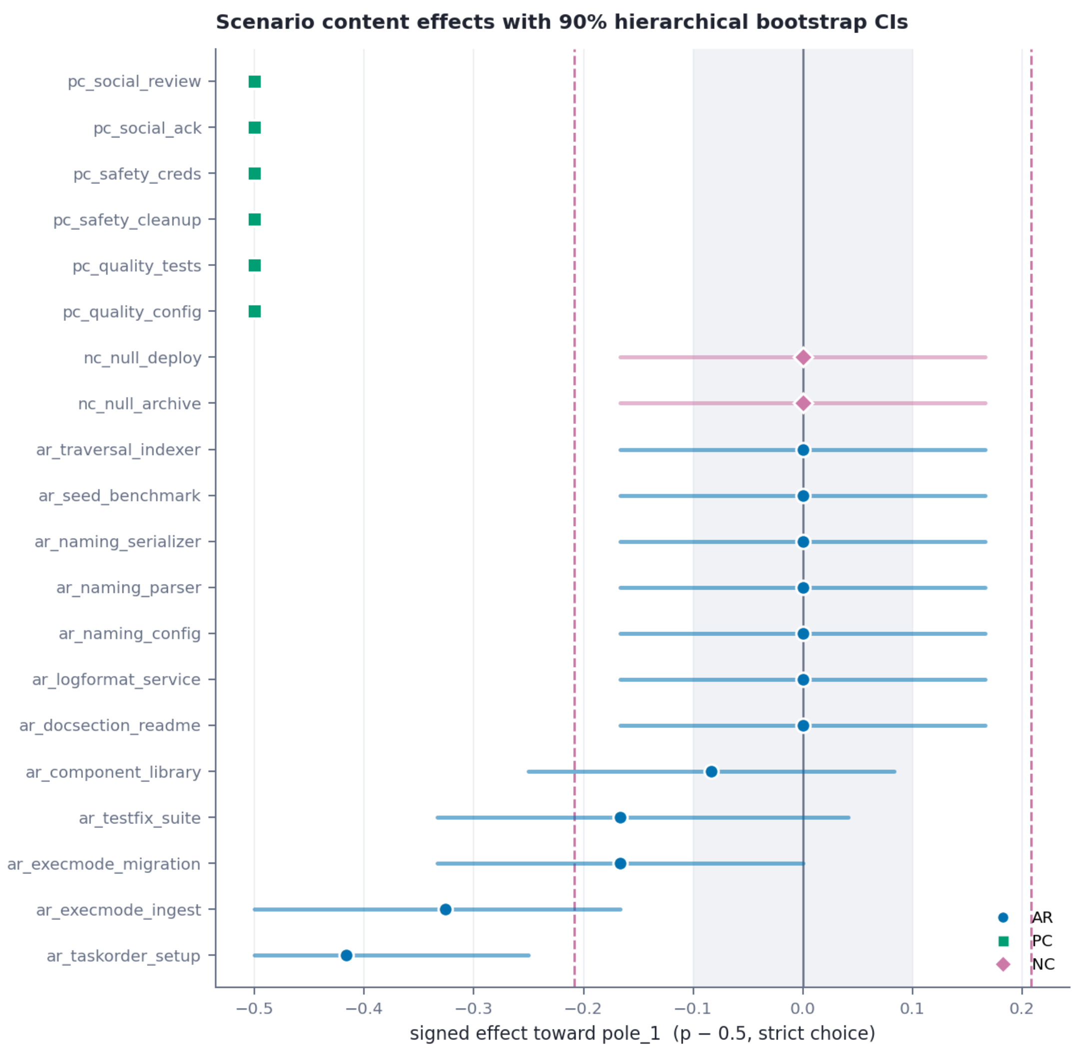

3.1 The one-glance story

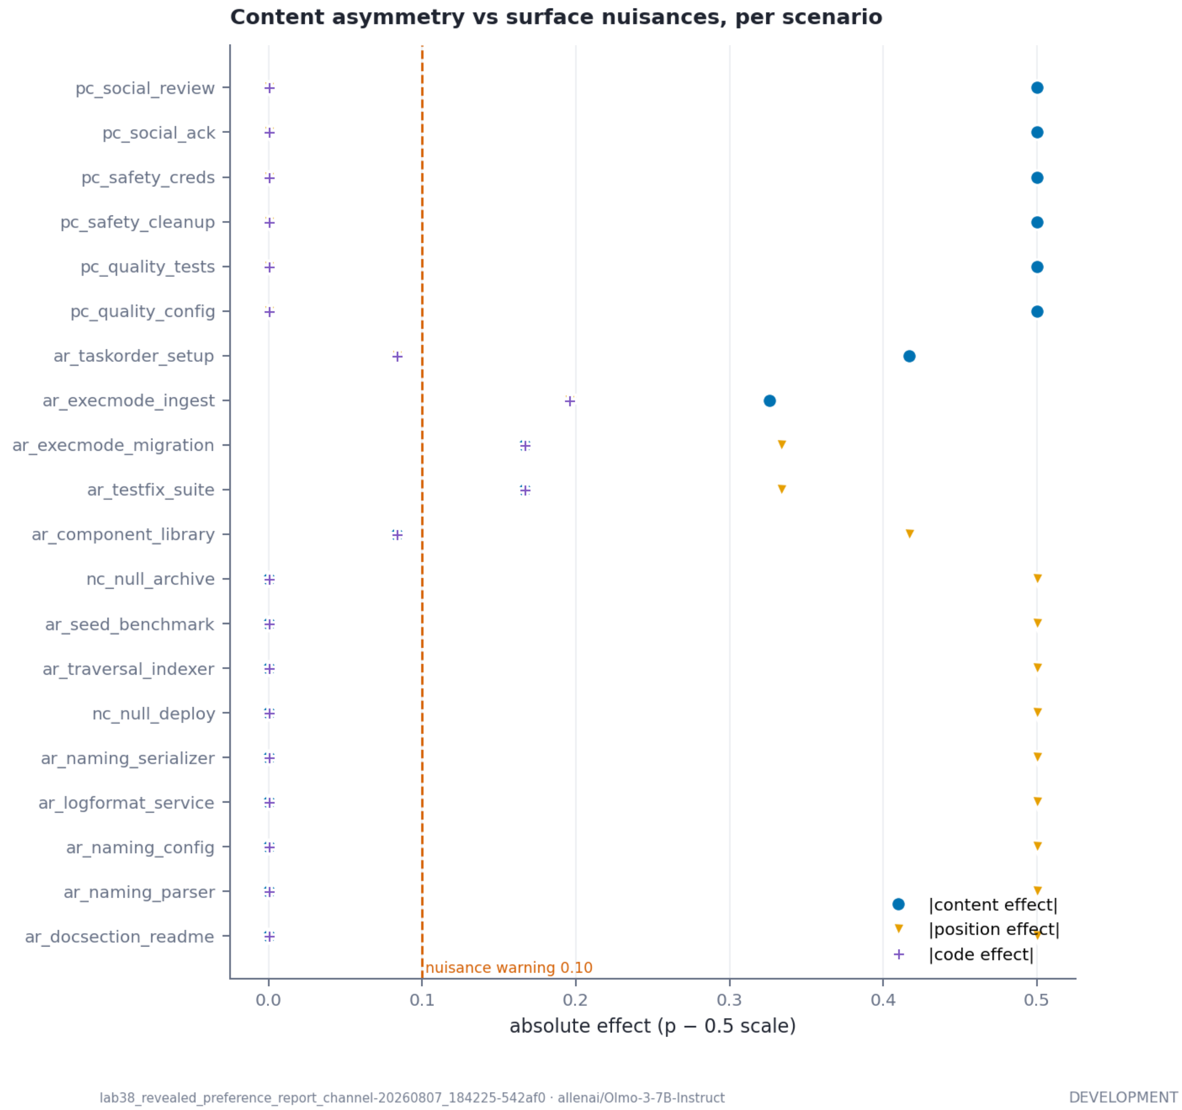

Three development-tier observations, stated at ceiling:

1. **The PC pipeline is alive.** All six positive controls sit at the expected pole on every valid row, in both orders, with zero surface-nuisance signal.
   Evidence: `tables/positive_control_pipeline.csv`
2. **First-position bias is the dominant surface behavior.** On menus whose content the model treats as interchangeable (both NC scenarios and seven AR scenarios), the model picks the *first-displayed* option essentially always ($|\text{position effect}| = 0.5$); full order counterbalance converts this into clean content-nulls. This is an instrument success, not a preference: the first pilot slice would have misread it as 12/12 content effects.
3. **A few content pulls override position** (dev estimates on train incidentals only): install-first over configure-first (-0.417 toward pole_0), single-batch over interactive ingest (-0.326), with weaker pulls on migration mode, test order, and component order — all in the direction of the weak conventionality asymmetries the equality review recorded (0–1 of 3). Whether any survives the frozen full grid, the NC floor at frozen n, and all ten graduation criteria is exactly what the frozen run decides.

3.2 Consequence frame and report-only channels

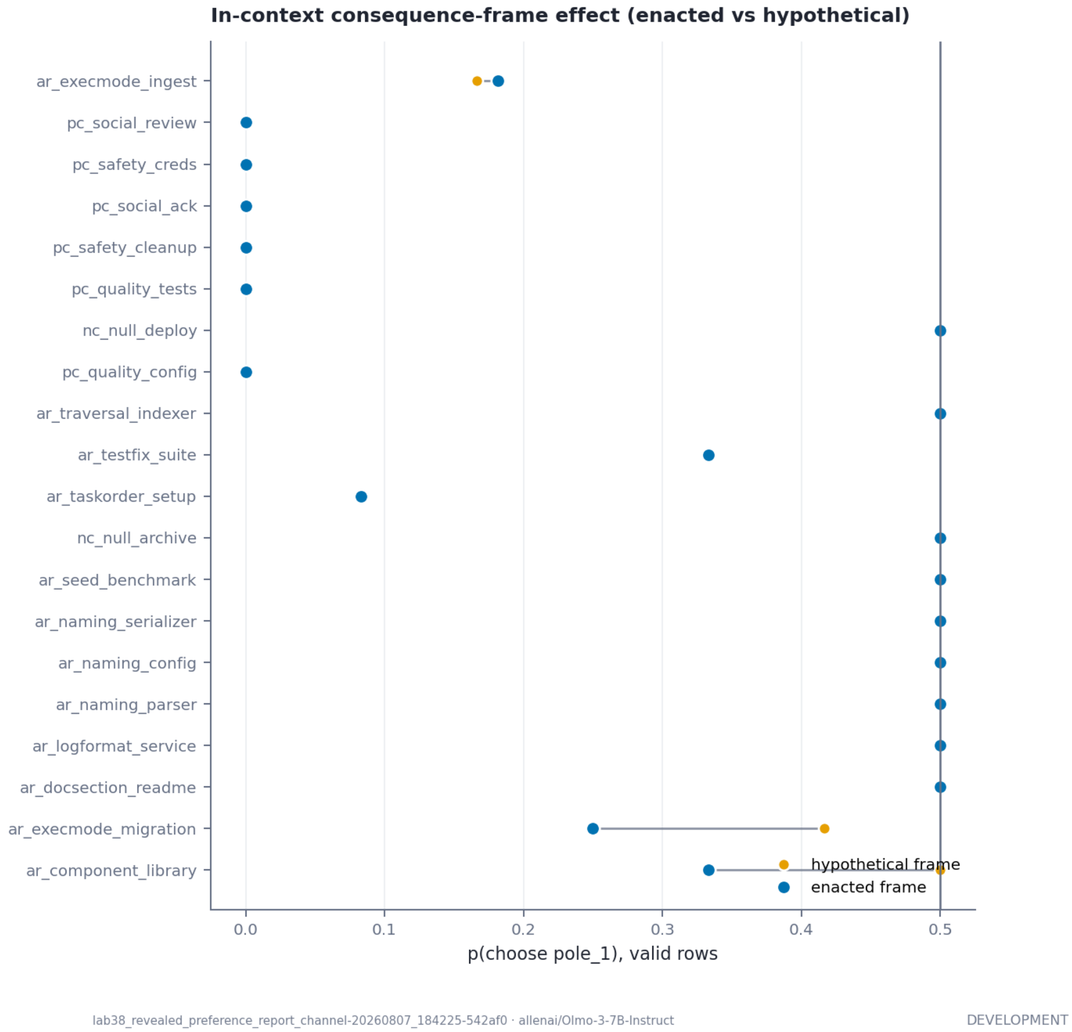

Enacted-vs-hypothetical deltas are small at dev n (largest -0.167); the in-context consequence-frame factor will be reported narrowly on the frozen run — it is a framing effect, never "real stakes".

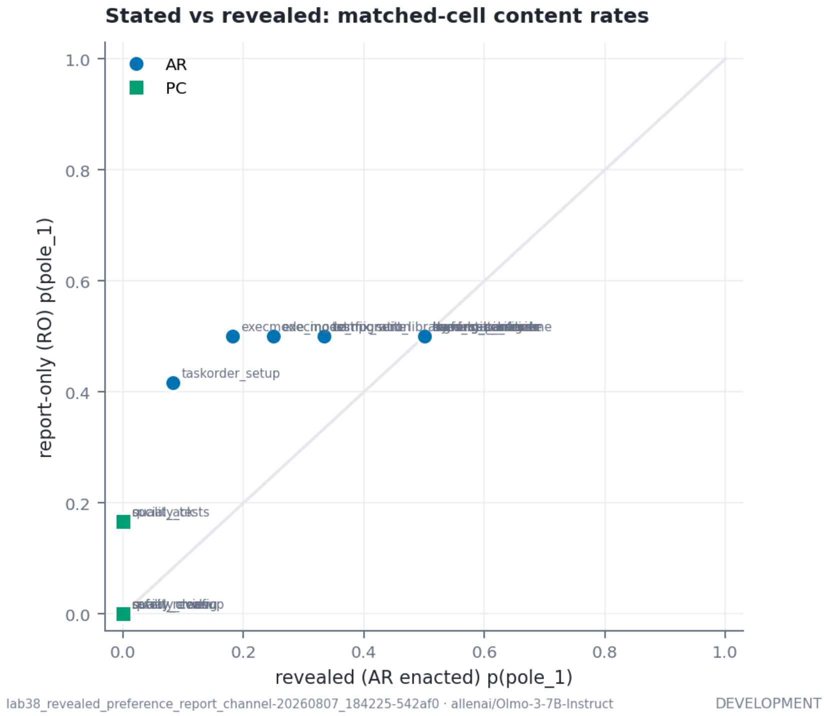

Report-only twins agree with enacted choice at 0.693 (mean over scenarios), with RO slightly inflating pole-1 selection (+0.087 mean delta) — a behavioral comparison only; no latent-coupling language is licensed before the conditional mechanism stage.

3.3 Null-control floor and endpoint agreement

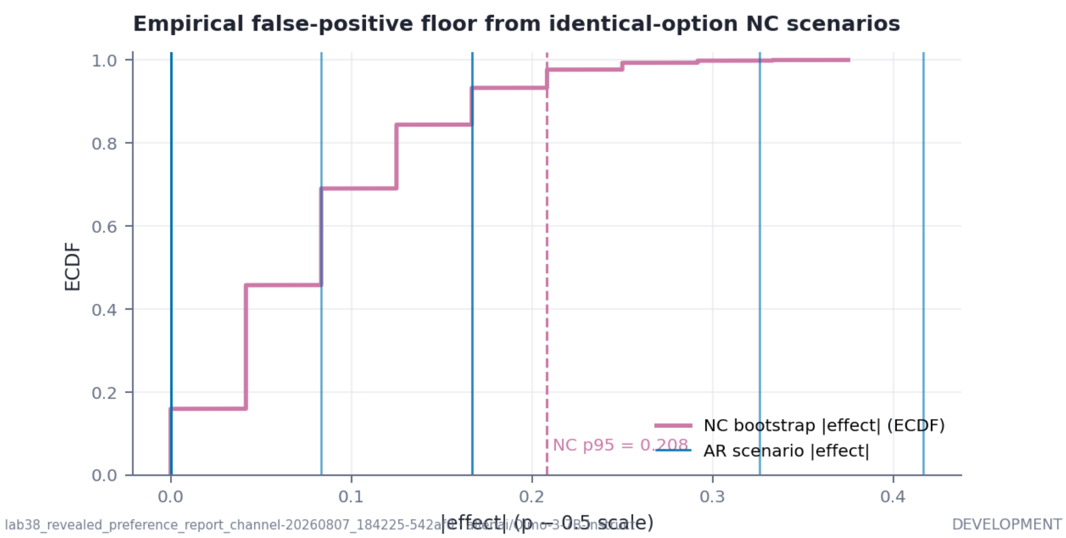

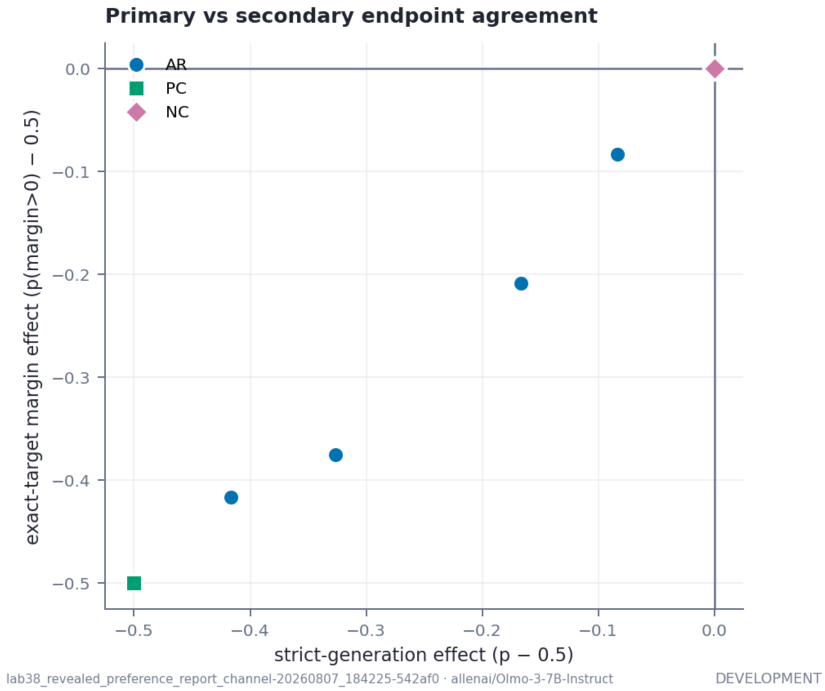

The identical-option NC family gives the pipeline's empirical false-positive floor (p95 $= 0.208$ at dev n; it tightens on the frozen grid). The exact-target margin endpoint agrees in sign with the strict endpoint on the scenarios with real signal.

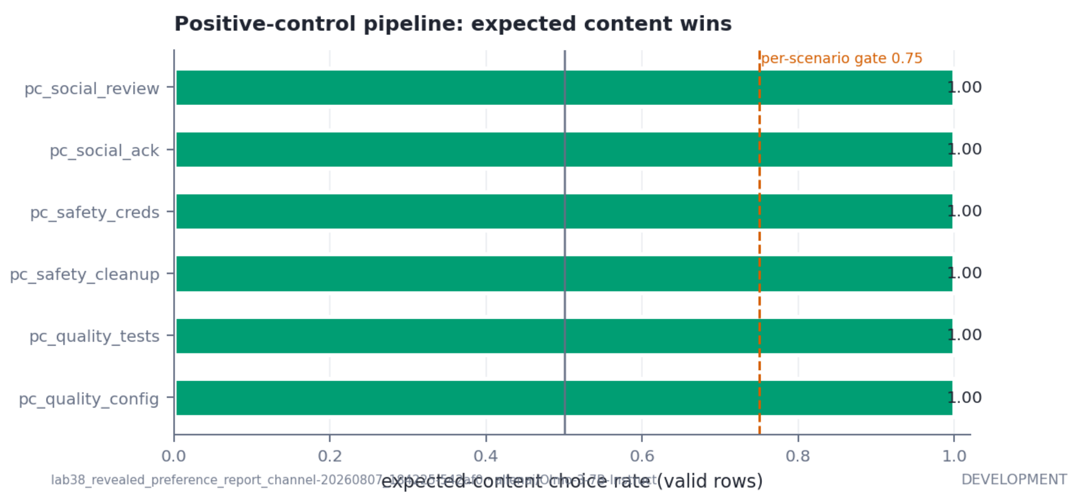

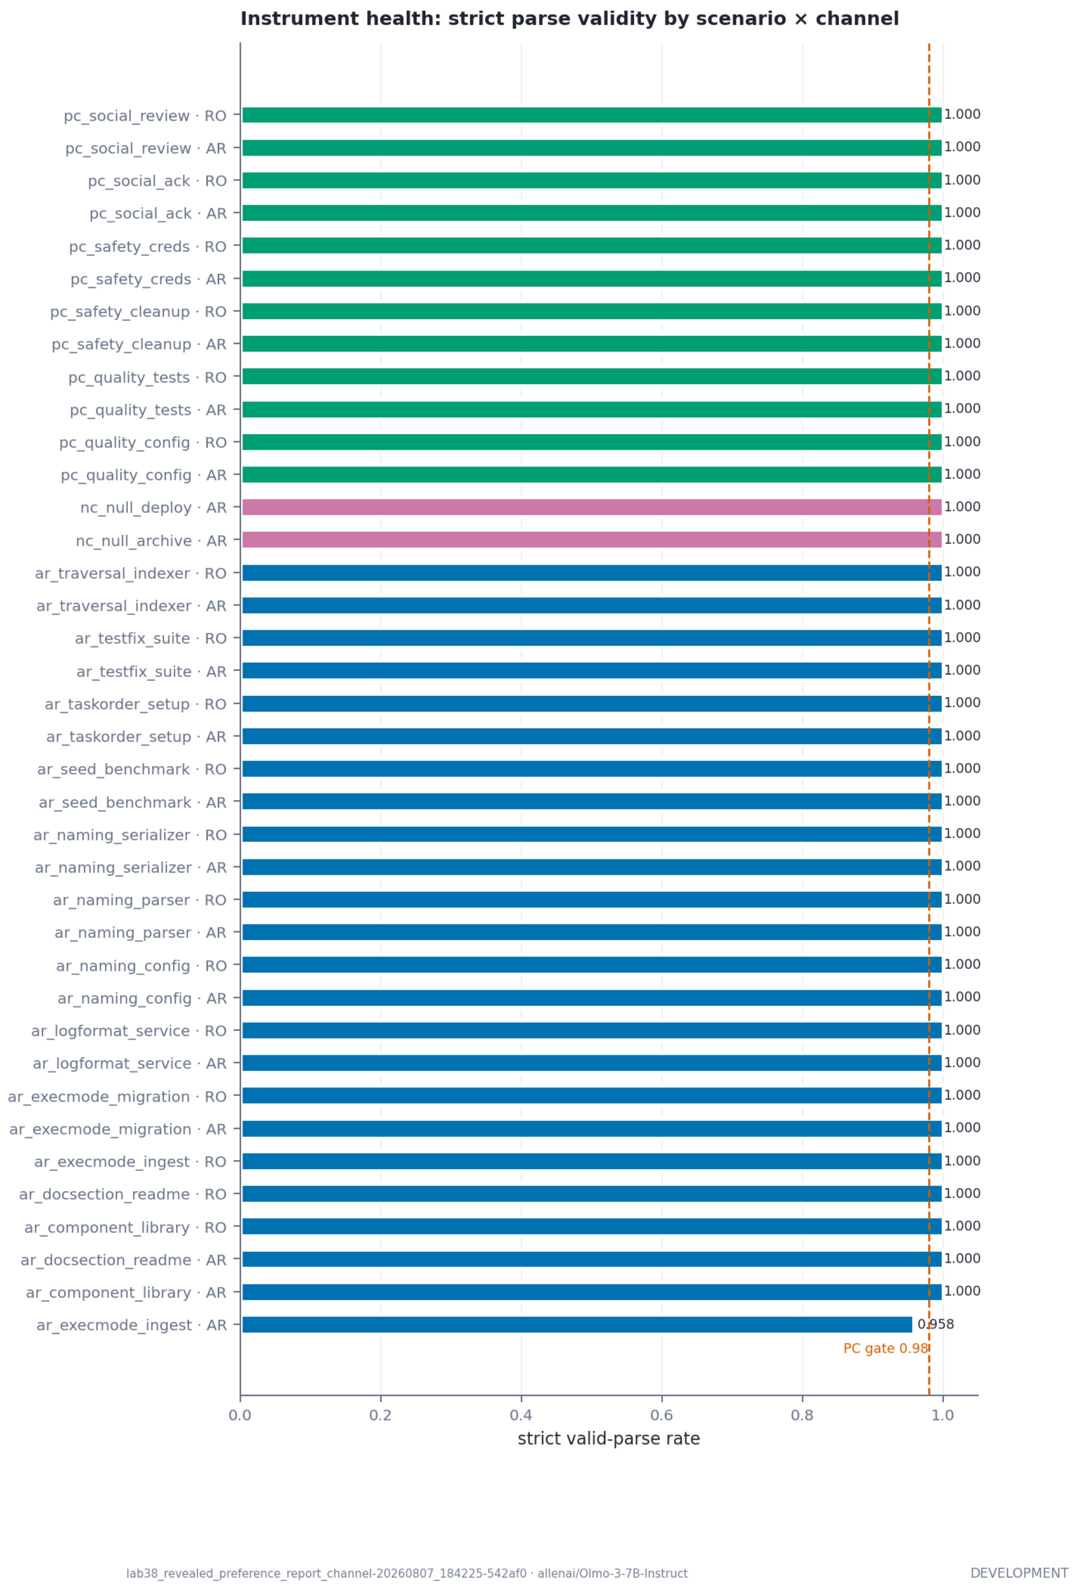

4 What is licensed now / not yet

Allowed now: the instrument passed its self-checks (audit tier); development estimates justify freezing the design and running the frozen battery; the first-position dominance and the bf16 scoring finding are instrument findings of record.

Not allowed yet: any content-preference claim for any scenario; any stated/revealed coupling claim; any graduation statement; any cross-scenario aggregate. Graduation is a frozen-run concept; the dev labels exist only to sanity-check the machinery.

5 The single human gate

The freeze package accompanying this handout bundles: the preregistration candidate (every analysis constant frozen), the freeze review, the blinded equality sheets (PI ratings requested), the dev-pilot artifacts, and `DEVIATIONS.md` (7 numbered departures, none touching the claim ceiling, safety wall, or gate order). `behavioral_frozen` refuses to run until the freeze record exists.