

Lab 38 Phase 1, part 1 — instrument build and development pilot

Stated vs revealed preference: counterbalanced action-binding choice on Olmo-3-7B-Instruct

Branch `interp_preference_phase1` — 2026-08-07 — freeze-review input, pre-registration candidate attached

Scope. This is a living **development only** report. Every number below is development-tier evidence produced before the freeze; nothing here is a preference result, and no frozen-partition outcome was opened. The claim ceiling applies verbatim: the campaign studies *functional* choice, report, and their coupling — never wants, welfare, suffering, consent, experience, moral patienthood, introspective truth, or a preference workspace. This handout is a freeze-review input, not a review signature, PI approval, freeze commit, or publication claim.

1 What was built

A governed campaign under `interpretability/preference/` (plan `preference_1.1.md` + binding addendum; both hash-pinned at intake). The two draft generators and the design transcript were lost before intake (`missing_at_intake`, PI-confirmed); the `lab38.v2.phase1` instrument was built from scratch with every plan-P0 repair native:

- **Bank:** 2,320 rows — 12 arbitrary (AR), 6 positive-control (PC: 2 quality / 2 social / 2 safety), 2 null-control (NC, verbatim-identical options) scenarios $\times$ 5 incidentals $\times$ full counterbalance (2 orders $\times$ 2 display-label sets $\times$ 2 response-code maps $\times$ 2 consequence frames), plus a report-only (RO) twin for every AR/PC cell. Splits 3 train / 1 validation / 1 holdout per scenario; construct axes `naming_convention` $\times$ 3 and `execution_mode` $\times$ 2 make construct-level claims reachable. Scientific identity: every behavior-relevant field is canonically hashed into the item id. Regeneration is byte-deterministic.

Evidence: `pref1-schema-audit-v1`

- **Response-code contract:** opaque consonant-consonant-digit codes audited on the pinned tokenizer per the addendum-E procedure; AR pair `KP4/PK7` (neutral-prior gap 0.0031 nats), RO pair `VM2/GS2` (0.0275 nats); disjoint alphabets, distinct first tokens, no prefix relations; codes listed in display order (E12).

Evidence: `pref1-codebook-final-v1`

- **Strict parser:** exact-code-or-invalid, never guesses; 12/12 adversarial cases behave as specified. **Binding:** enacted valid choices execute exactly the parsed branch (4 scenarios carry model microtasks with deterministic validators); hypothetical, report-only, and invalid rows never execute (E10); wrong-branch executions are a hard stop.
- **Runner:** canonical-order batches, immutable per-item JSONL with fsync, atomic resume cursor that refuses config changes, interrupted-resume byte-parity proven, Drive mirroring at batch boundaries. **Equality review:** two blinded provisional agent passes (disagreements preserved), both pre-model; PI ratings gate the freeze. 46 unit + synthetic-analysis tests green.

Evidence: `pref1-gate-g2-instrument-v1`

2 Instrument finding: bf16 batched scoring is not batch-invariant

On the 7B in bf16, batched teacher-forced margins differed from single-row scoring by up to **0.250 nats** (kernel/shape numerics; strict generations were identical). Margins are therefore scored **single-row**: exact, batch-shape-free, and resume-invariant by construction. Per-run hard gates: single-row replay determinism and batched=single generations; the batched delta is recorded as an informational diagnostic. This is one of the most reusable numbers this part produced — any lab scoring small margins from batched bf16 forwards should check it.


Evidence: `diagnostics/batch_invariance.json`

3 Development pilot (696 rows, dev subset)

Dev subset = train incidentals $\times$ both orders $\times$ letter labels $\times$ both code maps $\times$ both frames. (The first pilot slice used one order only; NC immediately exposed the aliasing — a deterministic tie-break registered as a spurious $|\text{effect}| = 0.5$ — so the slice was widened before the freeze. Dev exists to catch exactly this.) Runtime: 10.7 rows/s measured; full frozen bank projects to ~ 9 GPU-minutes.

check	observed	gate
strict valid-parse rate (all 696)	0.9986	—
PC strict parse rate	1.000	≥ 0.98 PASS
PC expected-content rate	1.000	≥ 0.85 PASS
PC per-scenario minimum	1.000	≥ 0.75 PASS
binding execution (valid enacted)	1.000	$= 1.00$ PASS
wrong-branch executions	0	$= 0$ PASS
NC false-positive floor (p95, dev n)	0.208	informational
microtask follow-through pass	0.875	exploratory
AR $\leftrightarrow$ RO agreement (enacted)	0.693	descriptive


The single invalid generation in 696 rows was PK4 — a blend of the two AR codes; the strict parser correctly refused it and no branch executed.



lab38_revealed_preference_report_channel-20260807_184225-542af0 · allenai/Olmo-3-7B-Instruct

DEVELOPMENT


3.1 The one-glance story



Three development-tier observations, stated at ceiling:

1. **The PC pipeline is alive.** All six positive controls sit at the expected pole on every valid row, in both orders, with zero surface-nuisance signal.
Evidence: `tables/positive_control_pipeline.csv`
2. **First-position bias is the dominant surface behavior.** On menus whose content the model treats as interchangeable (both NC scenarios and seven AR scenarios), the model picks the *first-displayed* option essentially always ($|\text{position effect}| = 0.5$); full order counterbalance converts this into clean content-nulls. This is an instrument success, not a preference: the first pilot slice would have misread it as 12/12 content effects.
3. **A few content pulls override position** (dev estimates on train incidentals only): install-first over configure-first (-0.417 toward pole.0), single-batch over interactive ingest (-0.326), with weaker pulls on migration mode, test order, and component order — all in the direction of the weak conventionality asymmetries the equality review recorded (0–1 of 3). Whether any survives the frozen full grid, the NC floor at frozen n , and all ten graduation criteria is exactly what the frozen run decides.

3.2 Consequence frame and report-only channels



Enacted-vs-hypothetical deltas are small at dev n (largest -0.167); the in-context consequence-frame factor will be reported narrowly on the frozen run — it is a framing effect, never “real stakes”.


[illegible]

Report-only twins agree with enacted choice at 0.693 (mean over scenarios), with RO slightly inflating pole-1 selection (+0.087 mean delta) — a behavioral comparison only; no latent-coupling language is licensed before the conditional mechanism stage.


3.3 Null-control floor and endpoint agreement


The plot shows the Empirical Cumulative Distribution Function (ECDF) of the absolute effect size, scaled by the p=0.05 quantile, for two scenarios: NC bootstrap and AR scenario. The x-axis represents the scaled effect size, ranging from 0.0 to 0.4. The y-axis represents the ECDF, ranging from 0.0 to 1.0. The NC bootstrap ECDF (pink line) starts at 0.0 for x=0.0, jumps to approximately 0.15 at x=0.01, then to 0.45 at x=0.05, 0.7 at x=0.08, 0.85 at x=0.12, 0.93 at x=0.16, and reaches 1.0 at x=0.21. The AR scenario ECDF (blue line) is a step function that jumps from 0.0 to 1.0 at x=0.0. A vertical dashed pink line marks the NC p95 = 0.208. A legend in the bottom right corner identifies the two ECDFs.

effect (p=0.05 scale)	ECDF
0.0	0.0
0.01	0.15
0.05	0.45
0.08	0.70
0.12	0.85
0.16	0.93
0.21	1.00



The identical-option NC family gives the pipeline’s empirical false-positive floor ($p_{95} = 0.208$ at dev n ; it tightens on the frozen grid). The exact-target margin endpoint agrees in sign with the strict endpoint on the scenarios with real signal.





lab38_revealed_preference_report_channel-20260807_184225-542af0 · allenai/Olmo-3-7B-Instruct

DEVELOPMENT

4 What is licensed now / not yet

Allowed now: the instrument passed its self-checks (audit tier); development estimates justify freezing the design and running the frozen battery; the first-position dominance and the bf16 scoring finding are instrument findings of record.

Not allowed yet: any content-preference claim for any scenario; any stated/revealed coupling claim; any graduation statement; any cross-scenario aggregate. Graduation is a frozen-run concept; the dev labels exist only to sanity-check the machinery.

5 The single human gate

The freeze package accompanying this handout bundles: the preregistration candidate (every analysis constant frozen), the freeze review, the blinded equality sheets (PI ratings requested), the dev-pilot artifacts, and `DEVIATIONS.md` (7 numbered departures, none touching the claim ceiling, safety wall, or gate order). `behavioral_frozen` refuses to run until the freeze record exists.